

```
In [78]: import numpy as np
from sklearn.linear_model import LogisticRegression
from sklearn.datasets import load_iris
from sklearn.model_selection import train_test_split
import pickle as cp
import matplotlib.pyplot as plt
```

```

In [79]: class NBC:
    def __init__(self, feature_types, num_classes):
        self.feature_types = feature_types
        self.num_classes = num_classes

    def fit(self, X, y):
        classes, classCounts = np.unique(y, return_counts=True)

        pis = np.expand_dims(classCounts / y.shape[0], axis=1)

        allMus = []
        allSigmaSqr = []

        for i in range(classes.shape[0]):
            # Get class examples
            classExampleIdxs = np.argwhere(y==classes[i])[:,0]

            # Calculate class parameters
            # Empirical mean / Bernoulli distribution parameter
            mus = np.average(X[classExampleIdxs,:], axis=0)
            # Empirical variance
            sigmaSqr = np.var(X[classExampleIdxs,:], axis=0)

            allMus.append(mus)
            allSigmaSqr.append(sigmaSqr)

        allMus = np.array(allMus)
        allSigmaSqr = np.array(allSigmaSqr)

        self.pis = pis
        self.mus = allMus
        self.sigmaSqr = allSigmaSqr
        self.classes = classes
        self.classCounts = classCounts

    def calcRealProb(self, X, realIndices, clsIdx):
        mean = self.mus[clsIdx,realIndices]
        variances = self.sigmaSqr[clsIdx,realIndices]
        realXs = X[:,realIndices]
        probs = np.exp(-np.square(realXs-mean)/(2*(variances+1e-6)))/np.sqrt(2
*np.pi*(variances+1e-6))
        return probs

    def calcCatProb(self, X, binIndices, clsIdx):
        params = self.mus[clsIdx,binIndices]
        binXs = X[:,binIndices]
        probs = np.zeros_like(binXs)
        for i in range(params.shape[0]):
            feature = binXs[:,i]
            featureProbs = np.zeros_like(feature)
            featureProbs[feature == 1] = params[i]
            featureProbs[feature == 0] = 1 - params[i]
            probs[:,i] = featureProbs
        return probs

    def calcClassProb(self, X, clsIdx):

```

```

featureTypes = self.feature_types
binIndices = [i for i, x in enumerate(featureTypes) if x == 'b']
realIndices = [i for i, x in enumerate(featureTypes) if x == 'r']

realProbs = self.calcRealProb(X, realIndices, clsIdx)
catProbs = self.calcCatProb(X, binIndices, clsIdx)

# Ensure no zeros
realProbs[realProbs == 0] = 1e-6
catProbs[catProbs == 0] = 1e-6

realProbs = np.log(realProbs)
catProbs = np.log(catProbs)
realProbs = np.sum(realProbs,axis=1)
catProbs = np.sum(catProbs, axis=1)

return realProbs + catProbs + \
    (np.log(self.pis[clsIdx]))

def predict(self, X):
    classProbs = []
    for i in range(self.classes.shape[0]):
        classProbs.append(self.calcClassProb(X, i))
    classProbs = np.array(classProbs)
    predictedClassIdx = np.argmax(classProbs,axis=0)
    return self.classes[predictedClassIdx]

```

```

In [80]: nbc = NBC(feature_types=['b','r','b','r'], num_classes=2)
nbc.fit(np.array([[1],[0.5],[1],[0.5]],[[1],[0.5],[0],[0.5]],[[1],[0.5],[0],[
0.5]]]),np.array([[2],[0],[2]]))

```

```

In [81]: # Hand-in 1:
# The value of C: 0.5 * 10.0 = 5.0

```

```
In [82]: # Iris Dataset
iris = load_iris()
XIris, yIris = iris['data'], iris['target']
XTrainIris, XTestIris, yTrainIris, yTestIris = \
    train_test_split(XIris, yIris, test_size=0.25)
XTrainIrisExpand = np.expand_dims(XTrainIris,axis=2)
XTestIrisExpand = np.expand_dims(XTestIris,axis=2)
yTrainIrisExpand = np.expand_dims(yTrainIris,axis=1)
yTestIrisExpand = np.expand_dims(yTestIris,axis=1)

nbc = NBC(feature_types=['r','r','r','r'], num_classes=3)
nbc.fit(XTrainIrisExpand, yTrainIrisExpand)
nbcYTrainPredict = nbc.predict(XTrainIrisExpand)
nbcYTestPredict = nbc.predict(XTestIrisExpand)
nbcTrainAccuracy = np.mean(nbcYTrainPredict == yTrainIrisExpand)
nbcTestAccuracy = np.mean(nbcYTestPredict == yTestIrisExpand)

logReg = LogisticRegression(\
    solver='lbfgs', multi_class='multinomial', max_iter=1000)
logReg.fit(np.squeeze(XTrainIris), np.squeeze(yTrainIris))
logRegYTrainPredict = logReg.predict(np.squeeze(XTrainIris))
logRegYTestPredict = logReg.predict(np.squeeze(XTestIris))
logRegTrainAccuracy = np.mean(logRegYTrainPredict == yTrainIris)
logRegTestAccuracy = np.mean(logRegYTestPredict == yTestIris)

print("Naive Bayes Training Accuracy", nbcTrainAccuracy)
print("Naive Bayes Test Accuracy", nbcTestAccuracy)
print("Logistic Regression Training Accuracy", logRegTrainAccuracy)
print("Logistic Regression Test Accuracy", logRegTestAccuracy)
```

Naive Bayes Training Accuracy 0.9642857142857143

Naive Bayes Test Accuracy 0.9473684210526315

Logistic Regression Training Accuracy 0.9732142857142857

Logistic Regression Test Accuracy 0.9736842105263158

```

In [83]: # Congressional Voting Records
XVote, yVote = cp.load(open('voting.pickle', 'rb'))
featureTypes = ['b' for i in range(XVote.shape[1])]
XTrainVote, XTestVote, yTrainVote, yTestVote = \
    train_test_split(XVote, yVote, test_size=0.25)
XTrainVoteExpand = np.expand_dims(XTrainVote,axis=2)
XTestVoteExpand = np.expand_dims(XTestVote,axis=2)
yTrainVoteExpand = np.expand_dims(yTrainVote,axis=1)
yTestVoteExpand = np.expand_dims(yTestVote,axis=1)

nbc = NBC(feature_types=featureTypes, num_classes=2)
nbc.fit(XTrainVoteExpand, yTrainVoteExpand)
nbcYTrainPredict = nbc.predict(XTrainVoteExpand)
nbcYTestPredict = nbc.predict(XTestVoteExpand)
nbcTrainAccuracy = np.mean(nbcYTrainPredict == yTrainVoteExpand)
nbcTestAccuracy = np.mean(nbcYTestPredict == yTestVoteExpand)

logReg = LogisticRegression(\
    solver='lbfgs', max_iter=10000)
logReg.fit(XTrainVote, yTrainVote)
logRegYTrainPredict = logReg.predict(XTrainVote)
logRegYTestPredict = logReg.predict(XTestVote)
logRegTrainAccuracy = np.mean(logRegYTrainPredict == yTrainVote)
logRegTestAccuracy = np.mean(logRegYTestPredict == yTestVote)

print("Naive Bayes Training Accuracy", nbcTrainAccuracy)
print("Naive Bayes Test Accuracy", nbcTestAccuracy)
print("Logistic Regression Training Accuracy", logRegTrainAccuracy)
print("Logistic Regression Test Accuracy", logRegTestAccuracy)

```

Naive Bayes Training Accuracy 0.9482758620689655

Naive Bayes Test Accuracy 0.9310344827586207

Logistic Regression Training Accuracy 0.9770114942528736

Logistic Regression Test Accuracy 0.9137931034482759

```

In [84]: # Split data
def splitData(X,y):
    XTrain, XTest, yTrain, yTest = \
        train_test_split(X, y, test_size=0.25)
    XTrain = np.expand_dims(XTrain,axis=2)
    XTest = np.expand_dims(XTest,axis=2)
    yTrain = np.expand_dims(yTrain,axis=1)
    yTest = np.expand_dims(yTest,axis=1)
    return XTrain, XTest, yTrain, yTest

```

```

In [85]: # Learning Curves of NBC and LogReg
def getLearningCurveNBC(X, y, estimator, numPerm):
    dataPerc = np.linspace(0.1,1.0,10)
    trainAccuracies = []
    testAccuracies = []
    for i in range(len(dataPerc)):
        totTrainAccuracy = 0
        totTestAccuracy = 0
        for j in range(numPerm):
            XTrain, XTest, yTrain, yTest = splitData(X,y)
            numTrain = XTrain.shape[0]
            numTrainSub = int(dataPerc[i] * numTrain)
            XTrainSub = XTrain[:numTrainSub]
            yTrainSub = yTrain[:numTrainSub]
            estimator.fit(XTrainSub, yTrainSub)
            yTrainPredict = estimator.predict(XTrainSub)
            yTestPredict = estimator.predict(XTest)
            totTrainAccuracy += \
                np.mean(yTrainPredict \
                    == yTrainSub)
            totTestAccuracy += \
                np.mean(yTestPredict == yTest)
            trainAccuracies.append(totTrainAccuracy / numPerm)
            testAccuracies.append(totTestAccuracy / numPerm)
    return trainAccuracies, testAccuracies

def getLearningCurveLogReg(X, y, estimator, numPerm):
    dataPerc = np.linspace(0.1,1.0,10)
    trainAccuracies = []
    testAccuracies = []
    for i in range(len(dataPerc)):
        totTrainAccuracy = 0
        totTestAccuracy = 0
        for j in range(numPerm):
            XTrain, XTest, yTrain, yTest = train_test_split(X, y, test_size=0.
25)
                numTrain = XTrain.shape[0]
                numTrainSub = int(dataPerc[i] * numTrain)
                XTrainSub = XTrain[:numTrainSub]
                yTrainSub = yTrain[:numTrainSub]
                estimator.fit(XTrainSub, yTrainSub)
                yTrainPredict = estimator.predict(XTrainSub)
                yTestPredict = estimator.predict(XTest)
                totTrainAccuracy += \
                    np.mean(yTrainPredict \
                        == yTrainSub)
                totTestAccuracy += \
                    np.mean(yTestPredict == yTest)
                trainAccuracies.append(totTrainAccuracy / numPerm)
                testAccuracies.append(totTestAccuracy / numPerm)
    return trainAccuracies, testAccuracies

```

```

In [86]: # Hand-in 2
dataPerc = np.linspace(0.1,1.0,10)
nbc = NBC(feature_types=featureTypes, num_classes=2)
nbcVoteTrainCurve, nbcVoteTestCurve = \
    getLearningCurveNBC(XVote, yVote, nbc, 150)
nbcVoteTestCurve = np.array(nbcVoteTestCurve)
nbcVoteTestCurve = 1 - nbcVoteTestCurve

nbc = NBC(feature_types=['r','r','r','r'], num_classes=3)
nbcIrisTrainCurve, nbcIrisTestCurve = \
    getLearningCurveNBC(XIris, yIris, nbc, 150)
nbcIrisTestCurve = np.array(nbcIrisTestCurve)
nbcIrisTestCurve = 1 - nbcIrisTestCurve

logReg = LogisticRegression(max_iter=1000, solver='lbfgs')
logRegVoteTrainCurve, logRegVoteTestCurve = \
    getLearningCurveLogReg(XVote, yVote, logReg, 150)
logRegVoteTestCurve = np.array(logRegVoteTestCurve)
logRegVoteTestCurve = 1 - logRegVoteTestCurve


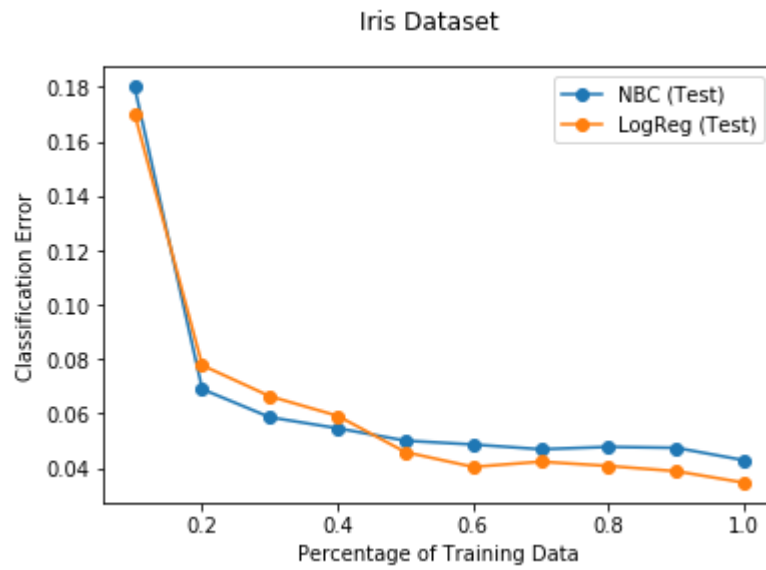
logReg = LogisticRegression(max_iter=1000, solver='lbfgs', multi_class='multinomial')
logRegIrisTrainCurve, logRegIrisTestCurve = \
    getLearningCurveLogReg(XIris, yIris, logReg, 150)
logRegIrisTestCurve = np.array(logRegIrisTestCurve)
logRegIrisTestCurve = 1 - logRegIrisTestCurve

plt.figure()
#plt.plot(dataPerc, nbcIrisTrainCurve, '-o', Label='NBC (Train)')
plt.plot(dataPerc, nbcIrisTestCurve, '-o', label='NBC (Test)')
#plt.plot(dataPerc, logRegIrisTrainCurve, '-o', Label='LogReg (Train)')
plt.plot(dataPerc, logRegIrisTestCurve, '-o', label='LogReg (Test)')
plt.xlabel("Percentage of Training Data")
plt.ylabel("Classification Error")
plt.suptitle("Iris Dataset")
plt.legend()

plt.figure()
#plt.plot(dataPerc, nbcVoteTrainCurve, '-o', Label='NBC (Train)')
plt.plot(dataPerc, nbcVoteTestCurve, '-o', label='NBC (Test)')
#plt.plot(dataPerc, logRegVoteTrainCurve, '-o', Label='LogReg (Train)')
plt.plot(dataPerc, logRegVoteTestCurve, '-o', label='LogReg (Test)')
plt.xlabel("Percentage of Training Data")
plt.ylabel("Classification Error")
plt.suptitle("Voting Dataset")
plt.legend()

```

Out[86]: <matplotlib.legend.Legend at 0x2195f5eca48>




```

In [87]: class NBC2:
    def __init__(self, feature_types, num_classes):
        self.feature_types = feature_types
        self.num_classes = num_classes

    def fit(self, X, y):
        classes, classCounts = np.unique(y, return_counts=True)

        pis = np.expand_dims(classCounts / y.shape[0], axis=1)

        allParams = []

        for i in range(classes.shape[0]):
            # Get class examples
            classExampleIdxs = np.argwhere(y==classes[i])[:,0]

            # Calculate class parameters
            # Bernoulli distribution parameter
            numOnes = np.count_nonzero(X == 1, axis=0)
            numZeros = np.count_nonzero(X == 0, axis=0)
            allParams.append(np.divide(numOnes, numOnes + numZeros))

        allParams = np.array(allParams)

        self.pis = pis
        self.params = allParams
        self.classes = classes
        self.classCounts = classCounts

    def calcCatProb(self, X, binIndices, clsIdx):
        params = self.params[clsIdx,binIndices]
        binXs = X[:,binIndices]
        probs = np.zeros_like(binXs)
        for i in range(params.shape[0]):
            feature = binXs[:,i]
            featureProbs = np.zeros_like(feature)
            featureProbs[feature == 1] = params[i]
            featureProbs[feature == 0] = 1 - params[i]
            probs[:,i] = featureProbs
        return probs

    def calcClassProb(self, X, clsIdx):
        featureTypes = self.feature_types
        binIndices = [i for i, x in enumerate(featureTypes) if x == 'b']
        realIndices = [i for i, x in enumerate(featureTypes) if x == 'r']

        catProbs = self.calcCatProb(X, binIndices, clsIdx)
        # Ensure no zeros
        catProbs[catProbs == 0] = 1e-6

        # Catch empty entries
        catProbs[catProbs != 2] = np.log(catProbs[catProbs != 2])
        catProbs[catProbs == 2] = 0

        # Calculate probability
        catProbs = np.sum(catProbs, axis=1)

```

```

        return catProbs + \
            (self.classCounts * np.log(self.pis[clsIdx]))

    def predict(self, X):
        classProbs = []
        for i in range(self.classes.shape[0]):
            classProbs.append(self.calcClassProb(X, i))
        classProbs = np.array(classProbs)
        predictedClassIdx = np.argmax(classProbs,axis=0)
        return self.classes[predictedClassIdx]

```

```

In [88]: # Congressional Voting Records (Full)
XVoteFull, yVoteFull = cp.load(open('voting-full.pickle', 'rb'))
featureTypes = ['b' for i in range(XVote.shape[1])]
nbc2 = NBC2(feature_types=featureTypes, num_classes=2)
nbcVoteFullTrainCurve, nbcVoteFullTestCurve = \
    getLearningCurveNBC(XVoteFull, yVoteFull, nbc2, 150)
nbcVoteFullTestCurve = np.array(nbcVoteFullTestCurve)
nbcVoteFullTestCurve = 1 - nbcVoteFullTestCurve

logReg = LogisticRegression(max_iter=1000, solver='lbfgs')
logRegVoteFullTrainCurve, logRegVoteFullTestCurve = \
    getLearningCurveLogReg(XVoteFull, yVoteFull, logReg, 150)
logRegVoteFullTestCurve = np.array(logRegVoteFullTestCurve)
logRegVoteFullTestCurve = 1 - logRegVoteFullTestCurve

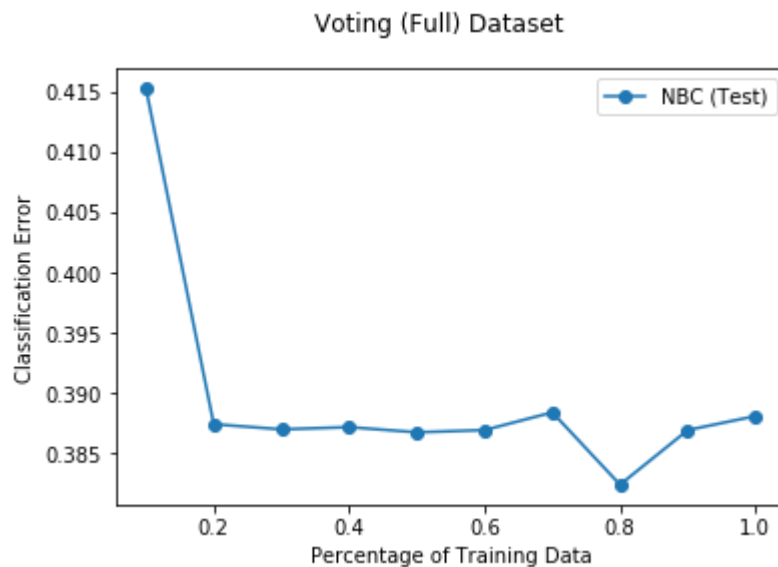

```

```

In [89]: dataPerc = np.linspace(0.1,1.0,10)
plt.figure()
plt.plot(dataPerc, logRegVoteFullTestCurve, '-o', label='LogReg (Test)')
plt.xlabel("Percentage of Training Data")
plt.ylabel("Classification Error")
plt.suptitle("Voting (Full) Dataset")
plt.legend()
plt.figure()
plt.plot(dataPerc, nbcVoteFullTestCurve, '-o', label='NBC (Test)')
plt.xlabel("Percentage of Training Data")
plt.ylabel("Classification Error")
plt.suptitle("Voting (Full) Dataset")
plt.legend()

```

Out[89]: <matplotlib.legend.Legend at 0x2195f69a188>



```
In [90]: # Congressional Voting Records (Full)
XVoteFull, yVoteFull = cp.load(open('voting-full.pickle', 'rb'))
# Replace missing features with mean
for i in range(XVoteFull.shape[1]):
    feature = XVoteFull[:,i]
    numZeros = np.sum(feature == 0)
    numOnes = np.sum(feature == 1)
    mean = numOnes / (numZeros + numOnes)
    feature[feature == 2] = mean
    XVoteFull[:,i] = feature

featureTypes = ['b' for i in range(XVote.shape[1])]
nbc = NBC(feature_types=featureTypes, num_classes=2)
nbcVoteFullTrainCurve, nbcVoteFullTestCurve = \
    getLearningCurveNBC(XVoteFull, yVoteFull, nbc, 150)
nbcVoteFullTestCurve = np.array(nbcVoteFullTestCurve)
nbcVoteFullTestCurve = 1 - nbcVoteFullTestCurve

logReg = LogisticRegression(max_iter=1000, solver='lbfgs')
logRegVoteFullTrainCurve, logRegVoteFullTestCurve = \
    getLearningCurveLogReg(XVoteFull, yVoteFull, logReg, 150)
logRegVoteFullTestCurve = np.array(logRegVoteFullTestCurve)
logRegVoteFullTestCurve = 1 - logRegVoteFullTestCurve
```

```

In [91]: dataPerc = np.linspace(0.1,1.0,10)
plt.figure()
plt.plot(dataPerc, logRegVoteFullTestCurve, '-o', label='LogReg (Test)')
plt.xlabel("Percentage of Training Data")
plt.ylabel("Classification Error")
plt.suptitle("Voting (Full) Dataset")
plt.legend()
plt.figure()
plt.plot(dataPerc, nbcVoteFullTestCurve, '-o', label='NBC (Test)')
plt.xlabel("Percentage of Training Data")
plt.ylabel("Classification Error")
plt.suptitle("Voting (Full) Dataset")
plt.legend()

```

Out[91]: <matplotlib.legend.Legend at 0x2195f77a388>

